

Implementing Deque in JavaScript

Foreword

The Double-Ended Queue, or Deque, is a fundamental and versatile data structure that serves as a cornerstone for numerous algorithms and system designs. While JavaScript's native `Array` provides methods that appear to mimic Deque functionality, a naive implementation built upon it suffers from critical performance penalties that can render applications inefficient, especially when handling large datasets. This whitepaper details a highly efficient, production-ready strategy for implementing a Deque in JavaScript, leveraging a plain object and pointer-based architecture to circumvent the performance pitfalls of naive approaches and achieve $O(1)$ constant-time complexity for all core operations.

1.0 Introduction to the Deque Data Structure

A Deque (pronounced "deck") is a sequential collection of elements that generalizes the functionality of a standard Queue. Where a Queue operates on a strict "First-In, First-Out" (FIFO) principle, allowing additions only at the rear and removals only from the front, a Deque provides the flexibility to add, remove, and inspect elements at both ends of the collection. This strategic capability makes it an invaluable tool for managing ordered data in scenarios requiring efficient access to both the first and last elements, such as in job schedulers, sliding window algorithms, and browsing history management.

1.1 Defining the Deque

The Deque is formally defined by the following characteristics:

- **Double-Ended Queue:** Its full name clarifies its core capability—it behaves like a queue that is open for operations at both ends.
- **Dual-End Operations:** Unlike a standard Queue, which restricts additions to the tail (rear) and removals to the head (front), a Deque permits the addition and removal of elements from both its head and its tail.

- **Limited Access:** True to its queue-like nature, a Deque typically restricts element inspection to the head and tail. Direct access to or modification of elements in the middle of the collection is not part of its standard interface.

1.2 Core Operations

The essential operations that define the Deque data structure are summarized below:

Operation	Description
Adding to the Head	Inserts a new element at the front (head) of the Deque.
Adding to the Tail	Inserts a new element at the back (tail) of the Deque.
Removing from the Head	Removes and returns the element at the front (head) of the Deque.
Removing from the Tail	Removes and returns the element at the back (tail) of the Deque.
Peeking at the Head	Returns the element at the head of the Deque without removing it.
Peeking at the Tail	Returns the element at the tail of the Deque without removing it.

Implementing these operations efficiently is the primary challenge, particularly within the context of JavaScript's built-in data structures.

2.0 The Inefficiency of an Array-Based Implementation

At first glance, the native JavaScript `Array` seems like an intuitive and convenient choice for implementing a Deque. It provides methods like `push()` and `pop()` for tail operations, and `unshift()` and `shift()` for head operations. However, this apparent convenience hides a significant performance bottleneck that makes arrays unsuitable for a high-performance Deque. This section will dissect those inefficiencies.

2.1 The Naive Approach

The straightforward method involves using a standard JavaScript `Array` to store the Deque's elements. `addTail` would map to `array.push()`, `removeTail` to `array.pop()`, `addHead` to `array.unshift()`, and `removeHead` to `array.shift()`. While tail operations are generally efficient, head operations introduce a severe performance penalty.

2.2 Analyzing the Performance Bottleneck

The critical performance issue with the array-based approach lies in how JavaScript arrays manage memory. When an element is added to the head of an array (at index `0`) using a method like `unshift()`, every other element in the array must be shifted one position to the right to make space.

This re-indexing is a **costly** and **time-consuming** operation. It requires an internal loop to move every subsequent element, from the first to the last. For an array containing n elements, this shifting process takes n steps. Consequently, adding or removing an element from the head is an **$O(n)$** operation. For large datasets, this linear time complexity is highly inefficient and leads to significant performance degradation.

This fundamental limitation necessitates an alternative architecture that can perform head operations without incurring the high cost of element shifting.

3.0 A High-Performance Architecture: The Object-Based Deque

The proposed solution is an architecture built not on an array, but on a plain JavaScript `Object` used as a map or key-value store, managed by two numeric pointers. This design fundamentally solves the shifting problem by abstracting element positions away from contiguous memory indices, allowing for constant-time insertions and deletions at both ends.

3.1 Core Components

The class structure for our high-performance Deque is built upon three key components:

- `items`: A plain JavaScript `Object` (`{}`) that serves as the key-value store for the Deque's elements. The keys will be numeric, representing positions, and the values will be the elements themselves.
- `head`: A numeric variable that acts as a pointer, storing the key of the first element in the `items` object. It is initialized to `0`.
- `tail`: A numeric variable that acts as a pointer to the *next available key* at the end of the Deque. It is also initialized to `0`. This design choice—having `tail` point to the next available position rather than the last element—is fundamental to the architecture's efficiency, simplifying the logic for both `addTail` and `peekTail` operations.

3.2 The Governing Principle

The central concept of this architecture is its use of an object's flexible keying system. Unlike array indices, which must be non-negative integers, object keys can be any number, including `0` and negative integers. This allows the `head` pointer to decrement "to the left" into negative keys when a new element is added to the front. This operation simply adds a new key-value pair to the object and updates a pointer, completely avoiding the need to shift any existing data.

The condition for determining if the Deque is empty is when the pointers are equal: `this.head === this.tail`.

This simple yet powerful architecture sets the stage for implementing all core Deque operations with maximum efficiency.

4.0 Deep Dive: Implementation of Core Operations

The following subsections provide a detailed breakdown of the logic for each primary Deque operation, demonstrating how the object-and-pointer architecture achieves superior performance.

4.1 Insertion Logic (`addHead` and `addTail`)

The logic for adding elements is distinct for the head and the tail, leveraging the pointer system to avoid data manipulation.

- `addHead` : To add an element to the front of the Deque, a simple two-step process is followed:
 1. The `head` pointer is decremented.
 2. The new element is assigned to the `items` object at the key indicated by the new `head` value.
- `addTail` : To add an element to the back of the Deque, a different two-step process is used:
 1. The new element is assigned to the `items` object at the key indicated by the *current* `tail` value.
 2. The `tail` pointer is incremented. This sequence works because the `tail` pointer is designed to always indicate the next open position, not the last filled one.

4.2 Deletion Logic (`removeHead` and `removeTail`)

Deletion logic first checks if the Deque is empty to avoid errors, returning `undefined` if so. Otherwise, it retrieves the element, removes it, and updates the appropriate pointer.

- `removeHead` : To remove the first element:
 1. Retrieve the element at the `head` key and store it for return.
 2. Use the `delete` keyword to remove the key-value pair from the `items` object, freeing memory.
 3. Increment the `head` pointer to designate the next element as the new head.
- `removeTail` : To remove the last element:
 1. Decrement the `tail` pointer to point to the last actual element.
 2. Retrieve the element at the new `tail` key and store it for return.
 3. Use the `delete` keyword to remove that key-value pair from the `items` object.

As a memory management optimization, if a deletion results in an empty Deque (`this.head === this.tail`), both pointers are reset to `0` to prevent them from drifting indefinitely into large positive or negative numbers over many operations.

4.3 Accessor and Utility Logic (`peek` , `size` , `isEmpty`)

Accessor and utility methods provide state information about the Deque without modifying it.

- `peekHead()` : Returns the item at `this.items[this.head]` .
- `peekTail()` : Returns the item at `this.items[this.tail - 1]` . The `1` is necessary because `tail` always points to the next available slot, so the last actual element is at the preceding key.
- `size()` : Returns the result of `this.tail - this.head` . This simple subtraction elegantly calculates the number of elements in the Deque. The mathematical distance between the pointers remains accurate regardless of whether `head` is positive, zero, or negative, demonstrating the robustness of this pointer-based approach.
- `isEmpty()` : Returns the boolean result of the comparison `this.head === this.tail` .

This operational logic directly translates into superior, constant-time performance for all core functions.

5.0 Performance Analysis and Justification

The justification for this object-based architecture is demonstrated conclusively by analyzing its time complexity against the naive array-based approach. The goal of any serious data structure implementation is to achieve the best possible performance, and this architecture delivers on that goal.

Operation	Array-Based (Naive)	Object-Based (Efficient)
<code>addHead</code> / <code>unshift</code>	$O(n)$	$O(1)$
<code>addTail</code> / <code>push</code>	$O(1)$	$O(1)$
<code>removeHead</code> / <code>shift</code>	$O(n)$	$O(1)$
<code>removeTail</code> / <code>pop</code>	$O(1)$	$O(1)$

The analysis is unequivocal: the object-based implementation provides **constant time, $O(1)$, performance for all primary operations**. This is achieved by completely eliminating the need for data shifting, which is the primary $O(n)$ bottleneck of the array-based methods for head operations. By simply manipulating pointers and adding or removing key-value pairs from an object, the cost of each operation remains constant, regardless of the number of elements in the Deque.

The following section presents the complete, production-ready code that encapsulates this high-performance design.

6.0 Complete Implementation and Demonstration

This section provides the full JavaScript `class` for the high-performance Deque, followed by a practical demonstration of its usage to illustrate its functionality and correctness.

6.1 The `Deque` Class in JavaScript

```
class Deque {  
  constructor() {  
    this.items = {};  
    this.head = 0;  
    this.tail = 0;  
  }  
}
```

```

}

isEmpty() {
  return this.head === this.tail;
}

addHead(element) {
  this.head--;
  this.items[this.head] = element;
}

addTail(element) {
  this.items[this.tail] = element;
  this.tail++;
}

removeHead() {
  if (this.isEmpty()) {
    return undefined;
  }
  const result = this.items[this.head];
  delete this.items[this.head];
  this.head++;

  if (this.isEmpty()) {
    this.head = 0;
    this.tail = 0;
  }
  return result;
}

removeTail() {
  if (this.isEmpty()) {
    return undefined;
  }
  this.tail--;
  const result = this.items[this.tail];
  delete this.items[this.tail];

```

```

    if (this.isEmpty()) {
      this.head = 0;
      this.tail = 0;
    }
    return result;
  }

  peekHead() {
    if (this.isEmpty()) {
      return undefined;
    }
    return this.items[this.head];
  }

  peekTail() {
    if (this.isEmpty()) {
      return undefined;
    }
    return this.items[this.tail - 1];
  }

  size() {
    return this.tail - this.head;
  }

  print() {
    if (this.isEmpty()) {
      console.log("Deque is empty");
      return;
    }
    let output = "";
    for (let i = this.head; i < this.tail; i++) {
      output += this.items[i] + (i < this.tail - 1 ? ' ↔ ': '');
    }
    console.log(output);
  }
}

```


6.2 Usage Demonstration

The following code demonstrates the instantiation and use of the `Deque` class, showing additions to both ends, followed by removals.

```
// 1. Create a new Deque instance
const d = new Deque();

// 2. Add several elements to the tail
d.addTail(10);
d.addTail(20);
d.addTail(30);

// 3. Add several elements to the head
d.addHead(40);
d.addHead(50);

// 4. Print the initial state of the Deque
console.log("Initial State:");
d.print();

// 5. Remove one element from the head and one from the tail
d.removeHead();
d.removeTail();

// 6. Print the final state of the Deque
console.log("\nFinal State:");
d.print();
```

Console Output:

```
Initial State:
50 ↔ 40 ↔ 10 ↔ 20 ↔ 30

Final State:
40 ↔ 10 ↔ 20
```

The output clearly illustrates the successful execution of the operations, maintaining the correct order of elements after additions and removals from

both ends.